
Distributed processing of JUNO datasets for improved energy reconstruction

Group: *Città Romanze*

- Pietro Cappelli
- Alberto Coppi
- Giacomo Franceschetto
- Nicolò Lai


Introduction



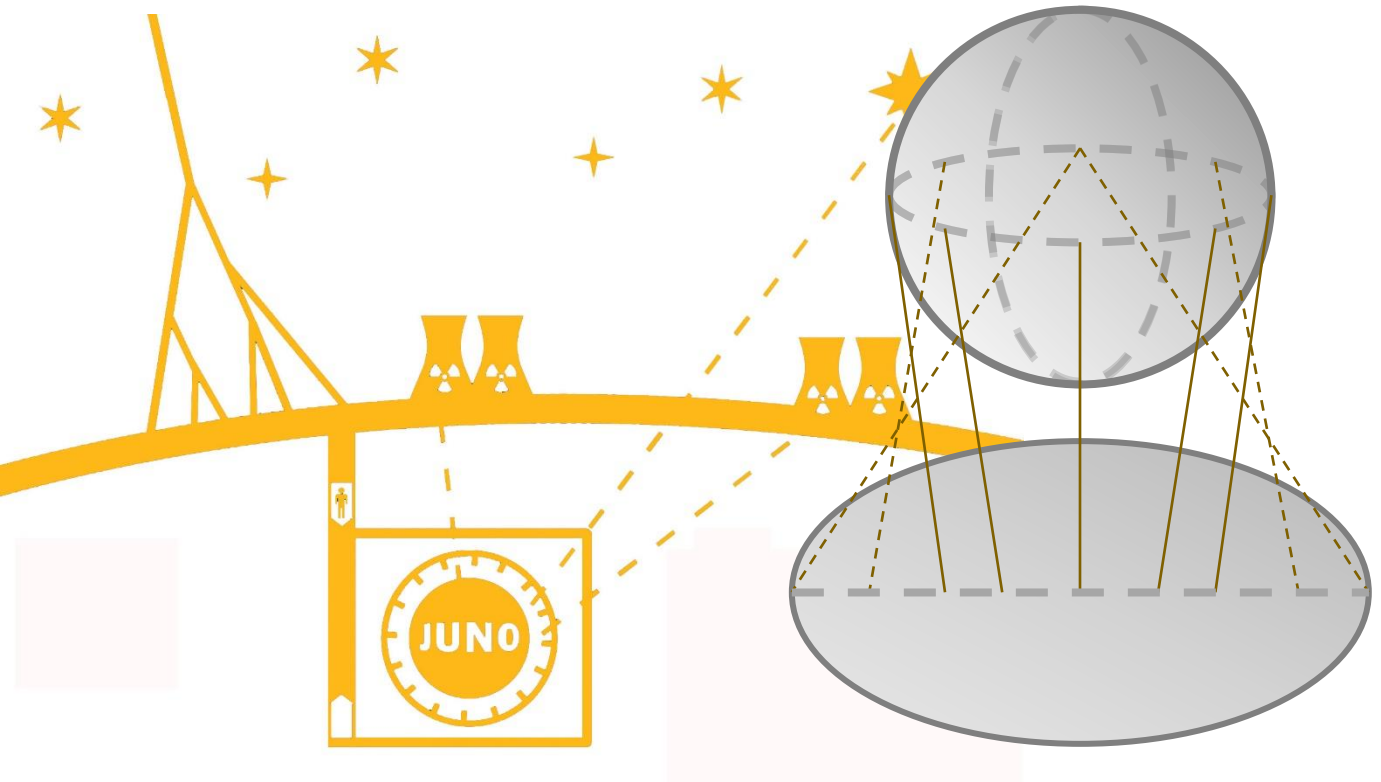
JUNO experiment & task overview



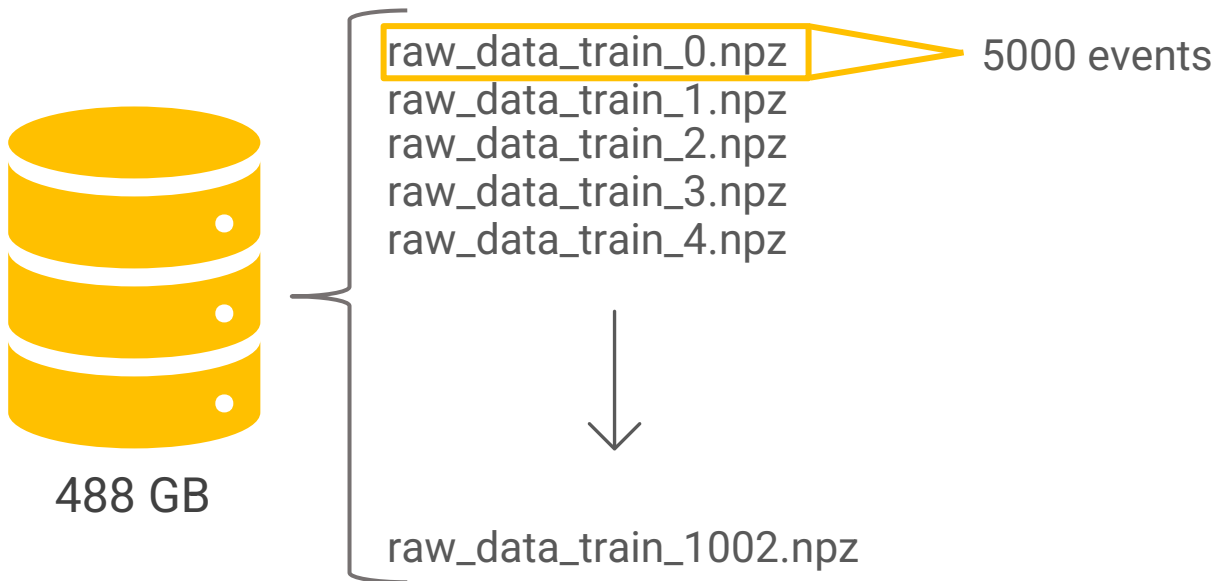
JUNO experiment & task overview



JUNO experiment & task overview



Dataset overview






Storage & Computing resources




Computing resources


We made use of **CloudVeneto** resources, specifically we had at our disposal:



Organized in **5 Virtual Machines**




And **one volume** holding data



5000 GB OF STORAGE



The Network File System





Motivations




Standard projection




Projection **pitfalls**



Rotation & Projection methods



Processing operations



Center of charge

$$x_c = \frac{\sum_i x_i \cdot c_i}{\sum_i c_i} \quad y_c = \frac{\sum_i y_i \cdot c_i}{\sum_i c_i} \quad z_c = \frac{\sum_i z_i \cdot c_i}{\sum_i c_i}$$


$$\varphi_C = \arctan\left(\frac{y}{x}\right) \quad \theta_C = \arctan \sqrt{\frac{x_c^2 + y_c^2}{z_c}}$$

Rotation

$$\phi_R = -\phi_c$$
$$\theta_R = -\theta_c + \frac{\pi}{2}$$

$$\begin{bmatrix} x' \\ y' \\ z' \end{bmatrix} = R_{yz}(\theta_R, \phi_R) \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

Processing operations



Center of charge

$$x_c = \frac{\sum_i x_i \cdot c_i}{\sum_i c_i} \quad y_c = \frac{\sum_i y_i \cdot c_i}{\sum_i c_i} \quad z_c = \frac{\sum_i z_i \cdot c_i}{\sum_i c_i}$$


$$\varphi_C = \arctan\left(\frac{y}{x}\right) \quad \theta_C = \arctan \sqrt{\frac{x_c^2 + y_c^2}{z_c}}$$

Rotation

$$\phi_R = -\phi_c$$
$$\theta_R = -\theta_c + \frac{\pi}{2}$$

$$\begin{bmatrix} x' \\ y' \\ z' \end{bmatrix} = R_{yz}(\theta_R, \phi_R) \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

Processing operations



Center of charge

$$x_c = \frac{\sum_i x_i \cdot c_i}{\sum_i c_i} \quad y_c = \frac{\sum_i y_i \cdot c_i}{\sum_i c_i} \quad z_c = \frac{\sum_i z_i \cdot c_i}{\sum_i c_i}$$

$$\varphi_C = \arctan\left(\frac{y}{x}\right) \quad \theta_C = \arctan \sqrt{\frac{x_c^2 + y_c^2}{z_c}}$$

Rotation

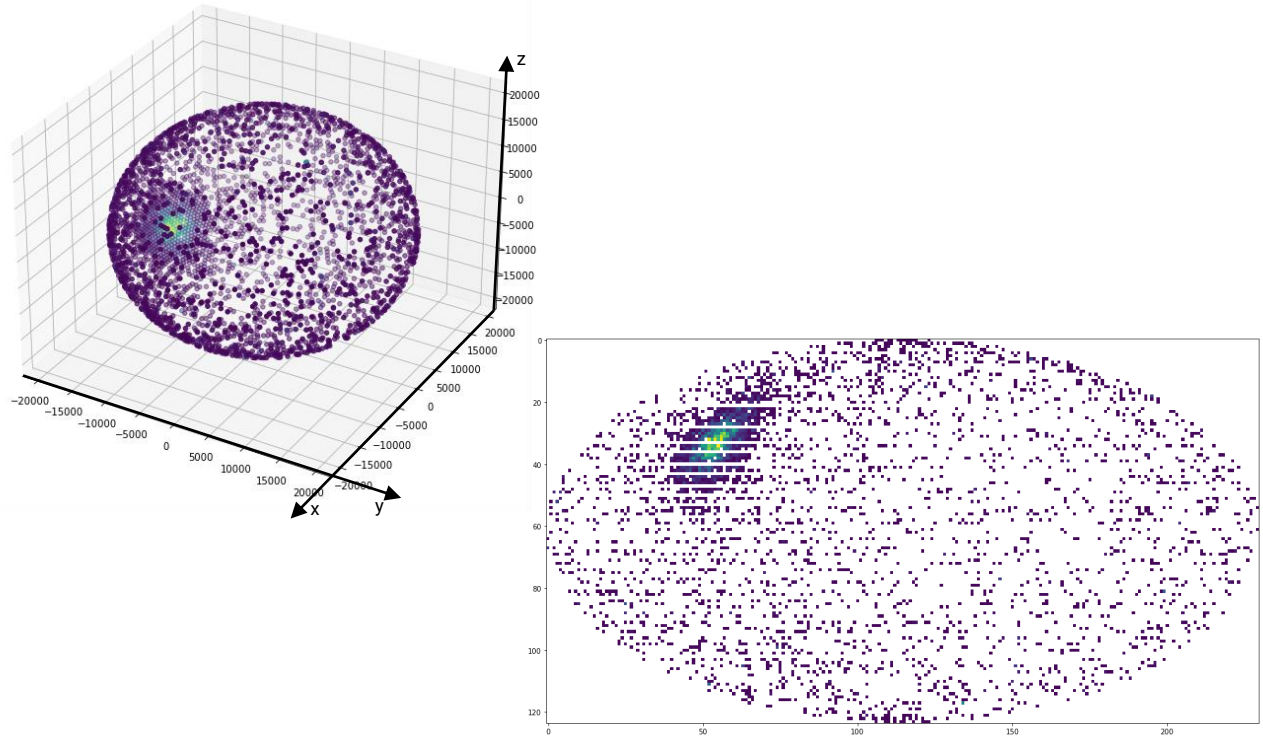
$$\phi_R = -\phi_c$$
$$\theta_R = -\theta_c + \frac{\pi}{2}$$

$$\begin{bmatrix} x' \\ y' \\ z' \end{bmatrix} = R_{yz}(\theta_R, \phi_R) \begin{bmatrix} x \\ y \\ z \end{bmatrix}$$

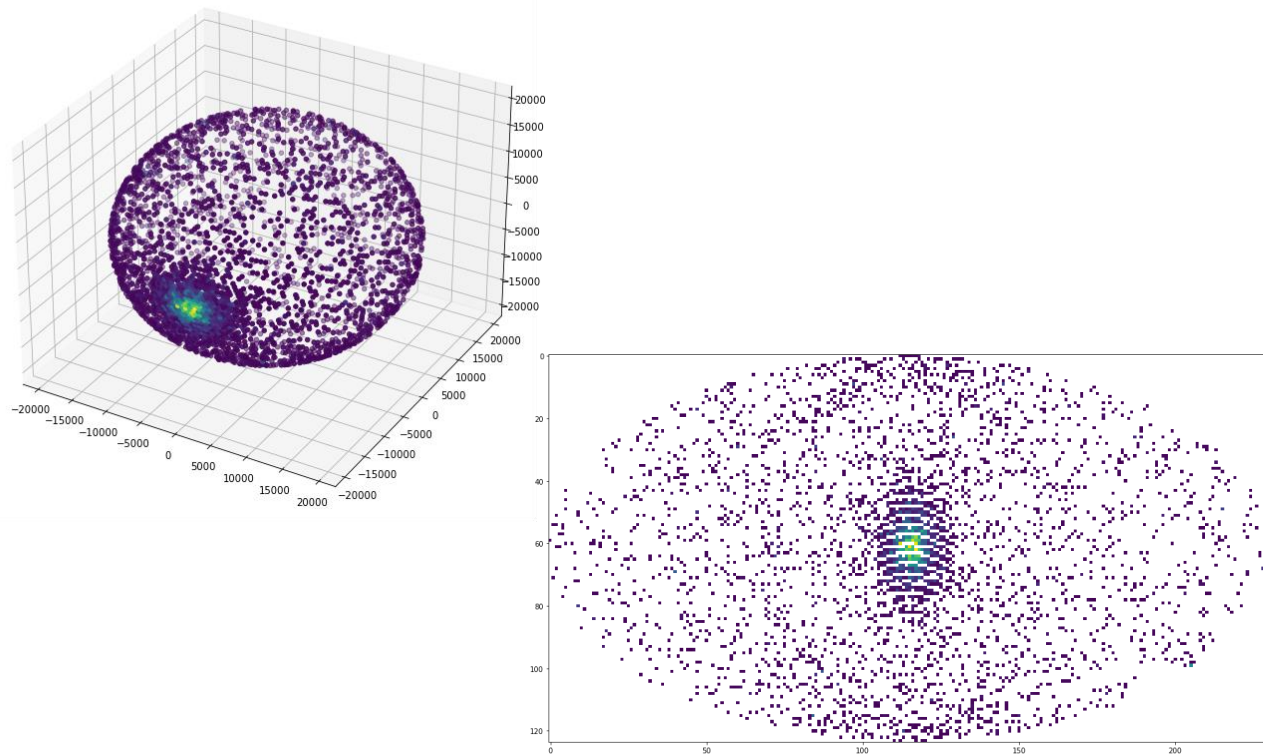
Projection


$$i_x = \left\lceil N_{eff} \cdot \frac{\tan^{-1}(x/y)}{\pi} \right\rceil + \frac{N_{max}}{2}$$

Standard mapping



Rotated event projection





Distributed algorithm implementation

Cluster setup



VM1

VM2



VM3



VM4



VM5



Cluster setup

```
cluster = SSHCluster(  
    ["10.67.22.39", "10.67.22.74", "10.67.22.27", "10.67.22.91", "10.67.22.60"],  
    connect_options = {"known_hosts": "/root/.ssh/known_hosts"},  
    worker_options   = {"nthreads": n_threads_wk, "n_workers": n_workers_vm},  
    scheduler_options = {"dashboard_address": ":8787"}  
)
```

SCHEDULER



4 workers
1 thread each



4 workers
1 thread each

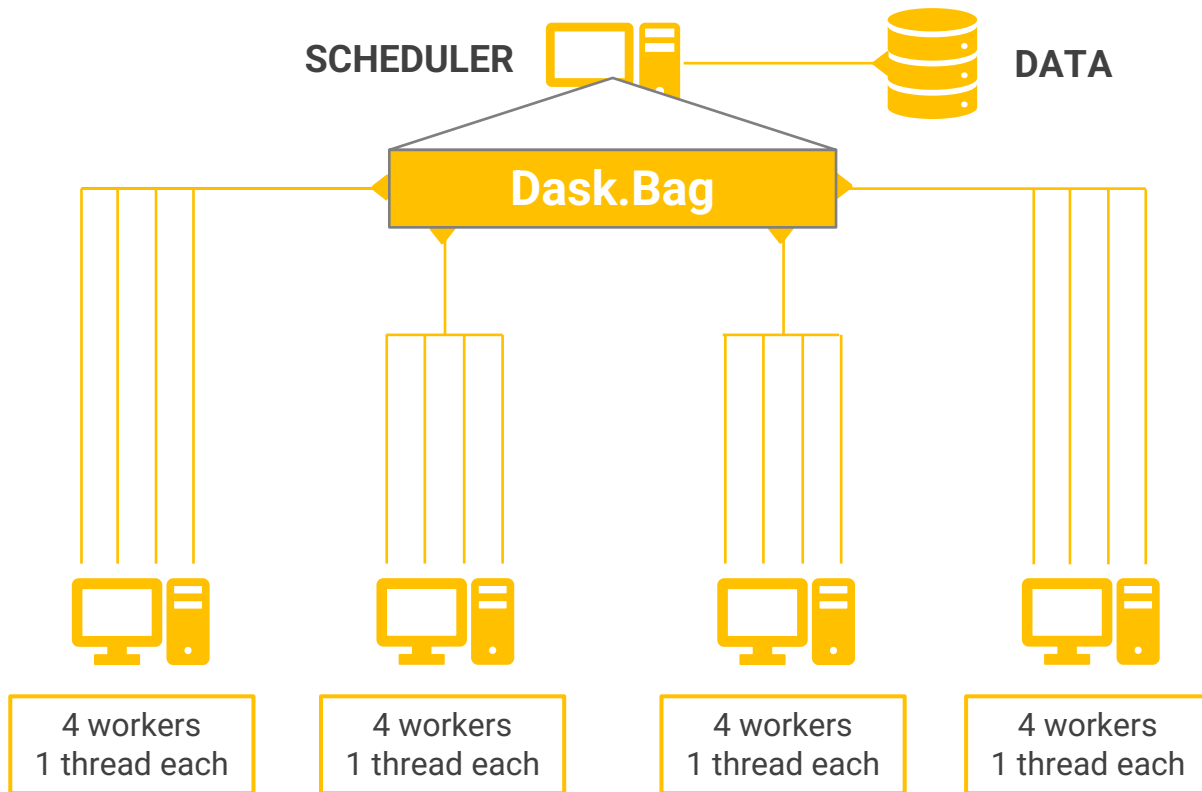


4 workers
1 thread each



4 workers
1 thread each

Data loading



The Distributed Processing

- *One-event* processing functions:

Rotation : rotate_ev()

Mapping : mapping_single_event()


- Processing distribution through

Dask.Bag.map()

```
rotated = db.map(rotate_ev, data_db)
mapped = db.map(mapping_single_event, rotated)
```


Lazy!

mapped.visualize()



Compute and Saving

```
images = mapped.compute()
```



Time performance benchmarks



The **three** parameters are

Workers



Number of workers to
spawn on a single virtual
machine

Threads




Size of the thread pool
within a single worker
process

Partitions




The granularity of data
parallelism


Benchmarks outline



Keeping the number of partitions fixed to sixteen




Setting the number of partitions equal to the total number of workers




Keeping the thread pool fixed to one single thread per worker

Benchmarks outline



Keeping the number of partitions fixed to sixteen



Setting the number of partitions equal to the total number of workers



Keeping the thread pool fixed to one single thread per worker




Workers




Threads

Keeping the number of partitions fixed to sixteen



Benchmarks outline



Keeping the number of partitions fixed to sixteen



Setting the number of partitions equal to the total number of workers




Keeping the thread pool fixed to one single thread per worker


Benchmarks outline



Keeping the number of partitions fixed to sixteen



Setting the number of partitions equal to the total number of workers



Keeping the thread pool fixed to one single thread per worker

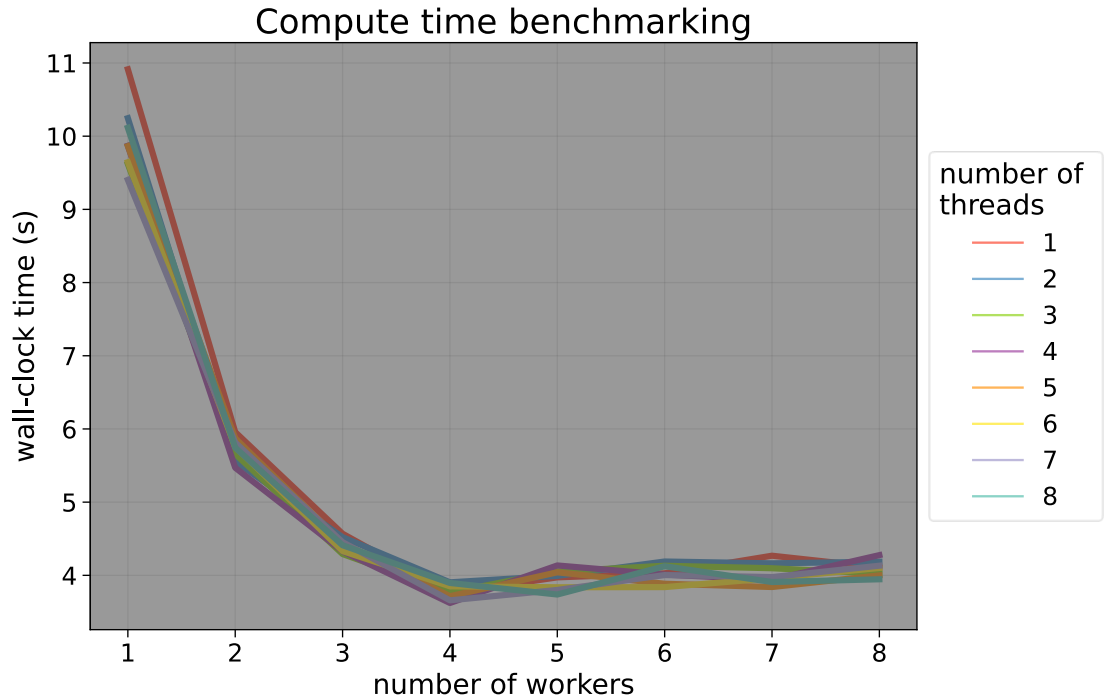


Workers



Threads


Setting the number of partitions equal to the total number of workers



Benchmarks outline



Keeping the number of partitions fixed to sixteen



Setting the number of partitions equal to the total number of workers




Keeping the thread pool fixed to one single thread per worker


Benchmarks outline



Keeping the number of partitions fixed to sixteen



Setting the number of partitions equal to the total number of workers



Keeping the thread pool fixed to one single thread per worker





Workers



Partitions


Keeping the thread pool fixed to one single thread per worker



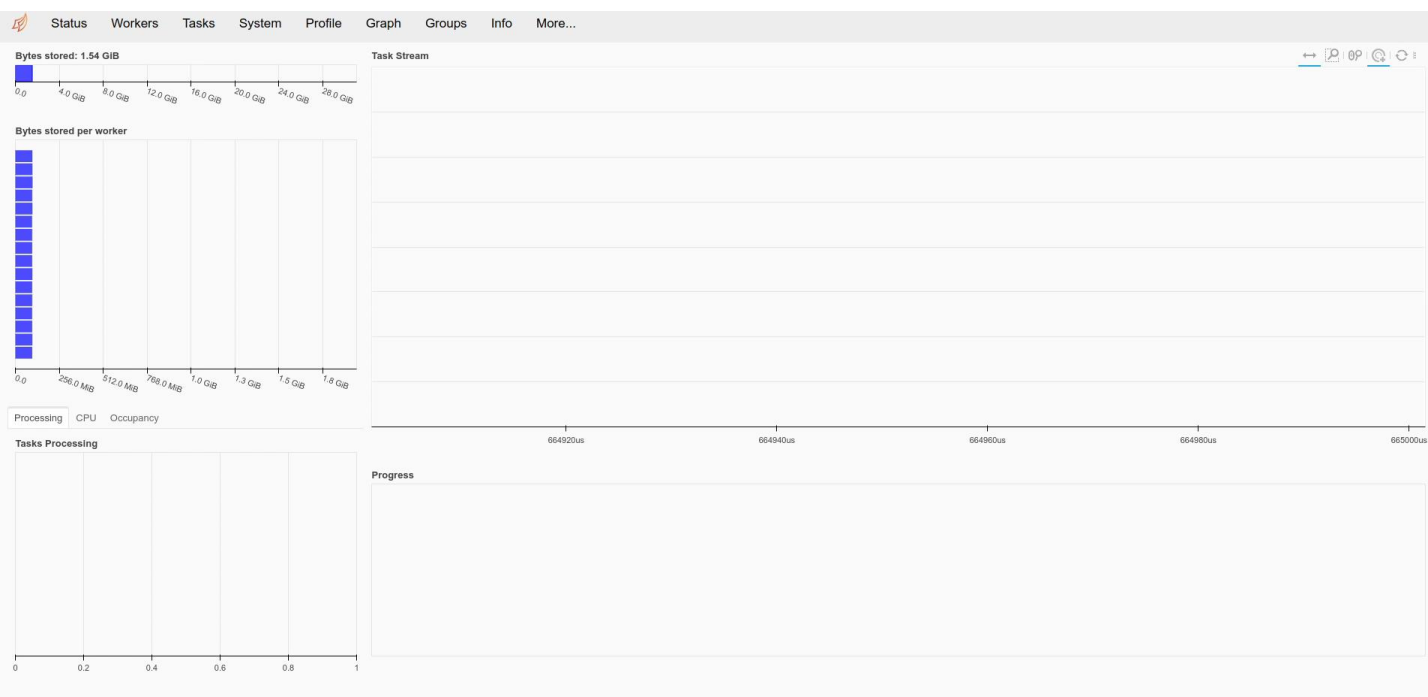


Distributed algorithm lazy implementation


Eager implementation pitfalls



Eager implementation pitfalls



Lazy implementation cluster configuration



2 workers
2 threads **4GB**
RAM each



2 workers
2 threads **4GB**
RAM each



2 workers
2 threads **4GB**
RAM each



2 workers
2 threads **4GB**
RAM each

Lazy implementation gimmicks

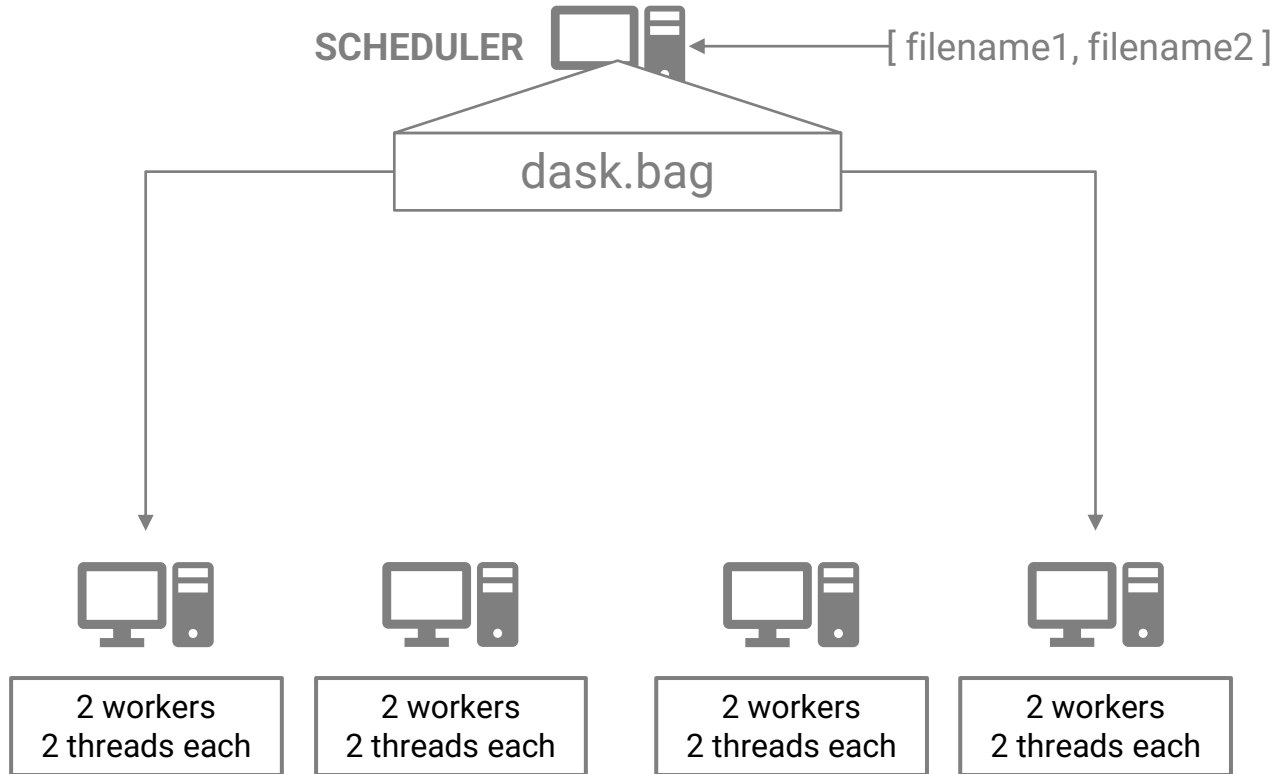
```
1 def load_bag(path, Nevents):
2     data_np = load(path)
3     data_mask = data_np[:, :Nevents]
4     bag = [np.vstack([ data_mask[j, i] for j in range(3)]) for i in range(data_mask.shape[1])]
5     del data_mask
6     del data_np
7     return bag
8 lazy_load_bag = dask.delayed(load_bag)
```

- Better RAM management
- delayed load_bag

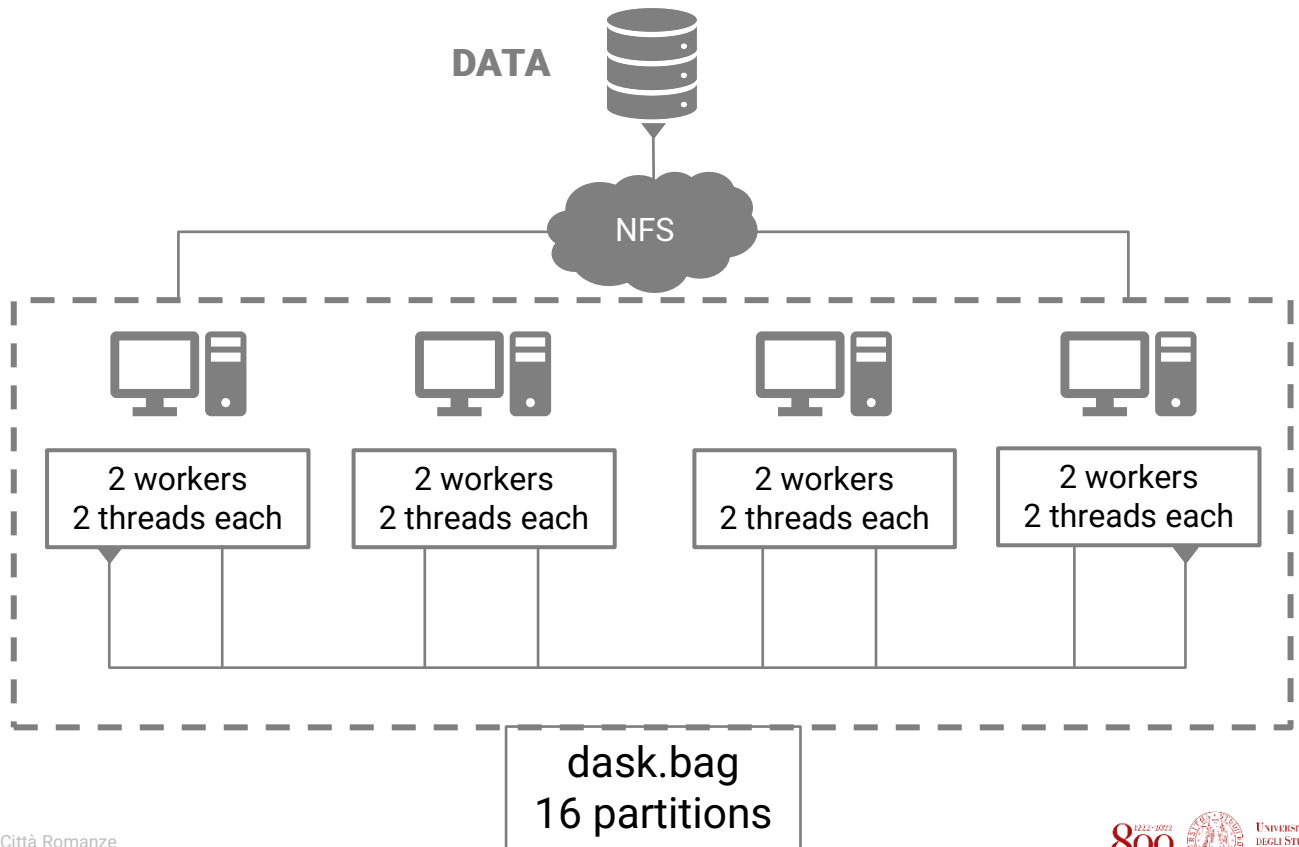
```
1 lazy_loaded = db.from_delayed([lazy_load_bag(data_folder + file_name, None)
2                                 for file_name in name_list[i:i+nfiles_cycle]])\
3     .repartition(n_workers_vm*n_threads_wk*4)
```

- Bag from delayed objects → Workers load files
- repartition over all available threads


Lazy implementation



Lazy implementation



Lazy implementation



**THANK YOU
FOR YOUR ATTENTION**


Città Romanze



QUESTIONS




Backup Slides



Lazy implementation gimmicks


```
1 future = client.compute(save, optimize_graph=False)
2 result = client.gather(future)
```

- Not optimized graph



Lazy implementation run

Task Stream




Lazy implementation run

Task Stream




Lazy implementation pitfalls




Lazy implementation pitfalls


Task Stream




Workers vs threads – 16 partitions





Workers vs threads – 16 partitions





Workers vs threads – 16 partitions





Workers vs threads – 16 partitions




Workers vs threads – 16 partitions




Workers vs threads – 16 partitions




Workers vs threads – partitions = workers





Workers vs threads – partitions = workers





Workers vs threads – partitions = workers





Workers vs threads – partitions = workers




Workers vs threads – partitions = workers




Workers vs threads – partitions = workers




Workers vs partitions – 1 thread





Workers vs partitions – 1 thread





Workers vs partitions – 1 thread



Workers vs partitions – 1 thread



Workers vs partitions – 1 thread



Workers vs partitions – 1 thread

